Entropy-Based Decoding for Retrieval-Augmented Large Language Models

Zexuan Qiu¹ Zijing Ou² Bin Wu³ Jingjing Li¹ Aiwei Liu⁴ Irwin King¹

¹The Chinese University of Hong Kong ²Imperial College London ³University College London ⁴Tsinghua University

Abstract

Augmenting Large Language Models (LLMs) with retrieved external knowledge has proven effective for improving the factual accuracy of generated responses. Despite their success, retrieval-augmented LLMs still face the distractibility issue, where the generated responses are negatively influenced by noise from both external and internal knowledge sources. In this paper, we introduce a novel, training-free decoding method guided by entropy considerations to mitigate this issue. Our approach utilizes entropy-based document-parallel ensemble decoding to prioritize low-entropy distributions from retrieved documents, thereby enhancing the extraction of relevant information of context. Additionally, it incorporates a contrastive decoding mechanism that contrasts the obtained low-entropy ensemble distribution with the high-entropy distribution derived from the model's internal knowledge across layers, which ensures a greater emphasis on reliable external information. Extensive experiments on open-domain question answering datasets demonstrate the superiority of our method.

1 Introduction

In recent years, Large language models (LLMs) have revolutionized natural language processing, showcasing remarkable performance across various downstream tasks (Brown et al., 2020; Ouyang et al., 2022; Touvron et al., 2023). However, they still struggle with hallucination due to the inaccuracy of parametric memory (Bubeck et al., 2023) and inherently tend to produce outdated information (Kasai et al., 2024). In contrast, explicitly augmenting LLMs with retrieved external knowledge from reliable datastores (Lewis et al., 2020; Borgeaud et al., 2022) can enable LLMs to generate content that exhibits less deviation from the truth, and benefit downstream knowledge-intensive tasks (Petroni et al., 2020).

Despite the success of retrieval-augmented LLMs, the augmented generation is still suboptimal due to the distractibility issue, where the generated responses are easily negatively affected by noise from both external knowledge and intrinsic model knowledge. As for the input context, LLMs' understanding of context can be explicitly distracted by irrelevant parts within the retrieved context (Shi et al., 2023a). A typical illustrative case is the "lost in the middle" distraction phenomenon observed in the synthetic multi-document question-answering scenario (Liu et al., 2024; Qiu et al., 2024), where the oracle document containing the correct answer is encircled by numerous retrieved distracting documents. In this scenario, LLMs frequently fail to deliver the correct answer unless the oracle document is strategically placed at the very beginning or end of the context. With regard to intrinsic knowledge, LLMs are easily implicitly distracted by the parametric knowledge acquired during pre-training. This is in conflict with retrieval-augmented generation which is expected to generate responses based on reliable retrieved context. Particularly in the domain of question answering, previous works (Longpre et al., 2021; Xie et al., 2023) show that LLMs stubbornly adhere to their built-in knowledge even when it conflicts with external knowledge.

How to eliminate the impact of the above-mentioned distractibility issue, so as to extract useful knowledge from the retrieved context for the input query, is our research focus. Although existing works strive to effectively leverage the retrieved context by directly fine-tuning retrieval-augmented LLMs (Lin et al., 2023; Zhang et al., 2024) or incorporating trainable encoder modules (Izacard and Grave, 2020; Yen et al., 2024), these approaches require additional training, rendering them potentially impractical in resource-constrained environments. In this paper, we propose a novel decoding method guided by entropy considerations to

simultaneously mitigate the impact of noisy information from both the external context and parametric knowledge. The proposed method can be seamlessly integrated into LLMs without requiring additional tuning.

Specifically, to enhance LLMs' ability to extract useful information from multiple retrieved documents, we let LLMs process each retrieved document in parallel and ensemble of the output distributions from each document to determine the next-token distribution, with the ensemble weights adaptively assigned based on the uncertainty of each document-conditioned distribution. At each generation step, documents with lower uncertainty (i.e., lower entropy) in the LLM output are given more attention during decoding. Ultimately, we obtain a low-entropy distribution aggregated among documents. Furthermore, to alleviate the potential distraction from parametric knowledge, we refine the next-token distribution by contrasting the obtained low-entropy distribution when feeding the retrieved documents, against the distribution without context. Here, we propose to use the distribution from the layer exhibiting the highest entropy without context for contrast, in order to highlight the proportional changes in token probabilities after introducing external knowledge.

The proposed decoding method shows an impressive performance in the synthetic challenging multi-document scenario (Liu et al., 2024) where the negative impact of retrieved distractor documents is emphasized. We further conduct extensive experiments across four LLMs of varying sizes on four diverse open-domain question answering tasks including NQ (Kwiatkowski et al., 2019), TriviaQA (Joshi et al., 2017), WebQ (Berant et al., 2013) and PopQA (Mallen et al., 2023). Experimental results confirm the superiority of our methods and validate the effectiveness of each component.

2 Methodology

In this work, we will investigate a new decoding approach for retrieval-augmented generation (RAG). Given an input query x, RAG first retrieves top-K relevant documents $D:=\{d_1,d_2,\ldots,d_K\}$ from the knowledge base via a retriever as external evidence, which is then incorporated with the query as the input to large language model parametrized by θ for generating a faithful response y.

Despite the integration of external knowledge, the decoding manner does matter to achieve desirable performance in RAG. A common approach, termed as "NAIVE RAG", involves concatenating the query x, the previously generated response $y_{< t}$, and the retrieved documents D as the input sequence, resulting in the following decoding method

$$p_{\theta}(y_t|D, x, y_{< t}) = p(y_t|d_0 \circ \cdots \circ d_k \circ x \circ y_{< t}), (1)$$

where \circ denotes the concatenation operation. Even though this method has demonstrated superior performance, Liu et al. (2024) highlight that it suffers from the "lost in the middle" distraction phenomenon, where LLMs tend to overlook the oracle document due to the distraction from other documents in D, unless the oracle document is placed at the very beginning or the end. Moreover, this simple approach does not consider the potential negative effects of the underlying parametric knowledge of LLMs. In the subsequent sections, we discuss how to mitigate these two issues simultaneously by entropy considerations.

2.1 Entropy-Based Document Ensemble

Instead of naively concatenating the documents D, we propose to alleviate the "loss in the middle" issue using the product-of-experts¹ ensemble approach (Hinton, 2002). Specifically, we model the log probability of the next-token distribution as

$$\log p_{\theta}(y_t|D) \propto \sum_{j=1}^{K} w_{j,t} \log p_{\theta}(y_t|d_j \circ x \circ y_{< t}), \quad (2)$$

where $\sum_j w_{j,t} = 1, \forall t, \ p_\theta(y_t|D,x,y_{< t})$ is denoted as $p_\theta(y_t|D)$ for short and $w_{j,t}$ denotes the weight of the j-th document on generating the token at the t-th time step. In Eq. (2), each document in D is concatenated with the query and the previously generated response. This combined input is then individually fed into the LLM. The output logit scores are subsequently averaged using the weights $w_{j,t}$. This ensemble approach, which leverages parallel decoding, helps mitigate position bias and provides a more effective means of utilizing the retrieved documents.

There are multiple choices to compute the weights. A straightforward option is to use uniform weighting, i.e., $w_{j,t} = \frac{1}{|D|}$. However, this method may fail to effectively extract valuable information when irrelevant documents are included among the top-K retrieved documents. Another

¹Empirically, we find that using product-of-experts and mixture-of-experts methods yield similar performance.




Figure 1: Overview of the decoding process of CLeHe.

option is to utilize a similarity score $s(d_j,x)$ between the query and the retrieved document, e.g., BM25, as the time-independent ensemble weights $w_{j,t} \propto s(d_j,x)$ (Lewis et al., 2020). However, this retriever-based scoring apporach may hinder the LLM's ability to extract relevant information when the retrieved documents contain numerous distractors, drastically reducing the factual accuracy of the responses as illustrated in § 4.1.

We posit that the uncertainty present in the next-token distribution inherently serves as a reliable indicator of the informativeness of the retrieved documents. Similar concepts have been employed in previous work to reduce hallucinations in LLMs (Van der Poel et al., 2022; Varshney et al., 2023). Consequently, we propose using an entropy-based score $w_{j,t}^{\rm H}$ as the preference weight for each document at each decoding step:

$$w_{j,t}^{H} = \frac{\exp^{(-H_{j,t})/\tau}}{\sum_{d_k \in D} \exp^{(-H_{j,t})/\tau}},$$

$$H_{j,t} = -\sum_{y_t \in \mathcal{V}} p_{\theta}(y_t|d_j) \log p_{\theta}(y_t|d_j),$$
(3)

where $p_{\theta}(y_t|d_j)$ denotes $p_{\theta}(y_t|d_j \circ x \circ y_{< t})$ for short, $\mathcal V$ represents the vocabulary set and τ is a hyperparameter controlling the concentration level of distributions. The motivation behind Eq. (3) is that the LLM can autonomously evaluate the significance of each document during the generation process. Intuitively, it implies that those document-conditioned distributions with lower uncertainty will be assigned higher weights. Such a time-dependent approach can effectively capture useful information from the retrieved documents at each generation step, thereby influencing the generation process more significantly. We refer to this method as LeEns (Low-entropy Ensemble).

2.2 Entropy-Based Contrastive Decoding

While **LeEns** can effectively help LLMs discern valuable evidence from external knowledge, the parametric knowledge of LLMs embedded during the pre-training phase might affect the answer generation, especially when these two types of knowledge conflict (Longpre et al., 2021; Xie et al., 2023). In this section, we propose to address this issue via entropy-based contrastive decoding.

Contrastive Decoding with PMI. Inspired by the success in contrastive decoding to mitigate hallucination of LLMs (Shi et al., 2023b; Li et al., 2022), we adjust the logit score z_t for the generated token $y_t \in \mathcal{V}$ at the t-th time step by incorporating the pointwise mutual information (PMI) between y_t and the document set D, given the query x:

$$z_{t} = \log p_{\theta}^{H}(y_{t}|D) + \beta \log \frac{p_{\theta}^{H}(y_{t}|D)}{p_{\theta}(y_{t}|x, y_{< t})}$$

$$= (1+\beta) \log p_{\theta}^{H}(y_{t}|D) - \beta \log p_{\theta}(y_{t}|x, y_{< t}),$$

$$(4)$$

where β is a positive coefficient to control the contrast intensity, and $p_{\theta}^{\mathrm{H}}(y_t|D)$ denotes the previously proposed entropy-based document ensemble distribution. Intuitively, PMI serves as a measurement of information gains. It is evident that the model tends to generate tokens with a high probability of $p_{\theta}^{\mathrm{H}}(y_t|D)$ and a low probability of $p_{\theta}(y_t|x,y_{< t})$. These tokens provide greater information gain for the next token generation. Consequently, incorporating PMI can enhance the model's reliance on external knowledge.

Layer-wise Contrast with High Entropy. To perform contrastive decoding, it is necessary to compute $p_{\theta}(y_t|x,y_{< t})$. This can be achieved by taking the hidden states from the last layer of LLMs

and passing them through the classification head. However, the distribution derived from the last layer may exhibit overconfidence, characterized by extremely low probabilities for most words and disproportionately high probabilities for a few. Such overconfidence can erroneously amplify external knowledge when conducting contrasting, potentially leading to false positive failures, as illustrated in Figure 7 in the Appendix.

To address this issue, we propose selecting the layer that contains the most "ambiguous" parametric knowledge among the layers as a proper reference for contrast. This allows the model to more effectively leverage external knowledge, reducing overconfidence and improving the accuracy of the generated outputs. Specifically, for a LLM consisting of a total of L layers, we denote the probability for $y_t \in \mathcal{V}$ in the l-th layer as:

$$p_{\theta}^{l}(y_{t}|x, y_{\leq t}) = \operatorname{softmax}(\boldsymbol{W}_{LM}h_{t-1}^{l}), \quad (5)$$

where h_{t-1}^l denotes the hidden state for layer l out of L, and $\mathbf{W}_{\rm LM}$ denotes the linear classification head in the LLM. At each decoding step, we dynamically select the layer with the maximum uncertainty for contrast:

$$l^* = \arg \max_{l \in \mathcal{L}} H_t^l,$$

$$H_t^l = -\sum_{y_t \in \mathcal{V}} p_{\theta}^l(y_t | x, y_{< t}) \log p_{\theta}^l(y_t | x, y_{< t}),$$
 (6)

where \mathcal{L} is the set of candidate layers, which are set as the last few layers of LLMs practically to ensure that each of them contains certain plausible information. Combining Eq. (3) and Eq. (6), the adjusted next-token distribution is formulated as:

$$y_t \sim \operatorname{softmax} \left[(1+\beta) \sum_{d_j \in D} w_{j,t}^{\mathsf{H}} \log p_{\theta}(y_t|d_j, x, y_{< t}) \right.$$

$$-\beta \log p_{\theta}^{l^*}(y_t|x, y_{< t}) \Big]. \tag{7}$$

Here, β represents the amplification intensity of external knowledge. When $\beta=0$, Eq. (7) degenerates to the proposed LeEns in § 2.1. The ultimate distribution in Eq. (7) can be interpreted as a two-stage ensemble process. Firstly, it ensembles the retrieved documents with uncertainty to generate a low-entropy distribution that more effectively captures the external knowledge within these documents. Secondly, it performs a contrastive ensemble by differentiating the logits of this low-entropy distribution from the high-entropy distribution of parametric knowledge selected across

different layers, thereby prioritizing factual information from external sources. In this regard, we term the method in Eq. (7) as **CLeHe** (Contrasting Low-entropy distribution with **H**igh-entropy distribution). Figure 1 illustrates the overall pipeline of the proposed CLeHe.

3 Related Works

Retrieval-Augmented Language Models. Enhancing large language models (LLMs) with information retrieved from external knowledge bases has proven effective for various knowledgeintensive tasks. Initially, mainstream research in retrieval-augmented language models (RALM) focused on leveraging retrieved knowledge during the pre-training phase of LLMs (Guu et al., 2020; Izacard et al., 2023; Borgeaud et al., 2022). To mitigate the computational costs, some studies have concentrated on lightweight fine-tuning methods to integrate retrieval capabilities into LLMs (Lewis et al., 2020; Lin et al., 2023; Zhang et al., 2024). Notably, models like FiD (Izacard and Grave, 2020) and CEPE (Yen et al., 2024) perform parallel encoding of multiple retrieved documents using a fine-tuned encoder, enabling decoder-only LLMs to more effectively capture and utilize external knowledge. Another approach leverages the in-context learning abilities of LLMs to incorporate external knowledge in a training-free manner (Ram et al., 2023; Shi et al., 2023c). The work most closely related to ours is REPLUG(Shi et al., 2023c), which utilizes the RAG-token model (Lewis et al., 2020) to perform parallel retrieval augmentation based on retrieval scores. However, we empirically demonstrate that focusing on the inherent uncertainty within the LLM's output distribution, rather than relying solely on pre-existing retrieval scores, can significantly improve the factual accuracy of content generated from retrieved documents.

Contrastive Decoding. The idea of contrastive decoding (CD) has been previously applied in controllable text generation to produce non-toxic by DExperts (Liu et al., 2021). Later, Li et al. (2022) formalized CD as a method to enhance open-ended text generation without any additional training by maximizing the difference in log probabilities between an expert LLM and an amateur LLM. This approach has demonstrated strong performance in various domains, including reasoning (O'Brien and Lewis, 2023) and neural machine translation (Waldendorf et al., 2024). CD can also

be interpreted as maximizing pointwise mutual information (PMI), which has proven effective in other contexts. For instance, Li et al. (2016) use a training objective that maximizes PMI to generate more diverse conversational responses, while CAD (Shi et al., 2023b) employs a PMI-adjust output probability distribution to resolve the knowledge conflict. Chuang et al. (2023) propose a decoding strategy that contrasts different layers of the same LLM to more effectively highlight factual knowledge obtained during pre-training. Similar principles are also applied in visual LLMs, where Leng et al. (2023) mitigate object hallucination by contrasting distributions derived from original and distorted visual inputs. Alternatively, our proposed CLeHe leverages layer-wise entropy-based contrastive decoding to prioritize external knowledge over the parametric knowledge inherent in the LLM itself.

4 Experiments

Baselines. We compare three training-free baselines. NAIVE: this method concatenates all retrieved documents directly along with the question to form the prompt for LLMs; REPLUG (Shi et al., 2023c): it utilizes a normalized retriever weight to ensemble during the decoding process; AvgEns: it follows the formulation in Eq.(2) and assigns the same weight to each document during each generation step. We verify the effectiveness of these decoding methods on 4 LLMs: LLAMA-2-7B, LLAMA-2-13B, Mistral-7B-v0.1, and LLAMA3-8B. Among them, LLAMA-3-8B supports an 8K context length, while the other three LLMs support a 4K context length. Few evaluation samples exceed 4k in context length, and they are removed to ensure the validity of the assessment.

Implementation Details. Our method introduces two hyperparameters: τ to control the relative importance of different documents during decoding; and β balancing contextual and parametric knowledge. We extract a subset from the WebQ training set for validation to determine the hyperparameter value for each LLM. Ultimately, τ is set as 0.25 for LLAMA-3-8B and as 0.1 for the other three models. For β , 5.0 is chosen for LLAMA-2-7B, while the other models are assigned a value of 0.25. During the inference, greedy decoding is utilized for reproducibility. When looking for the layer with the highest entropy, we focus our search exclusively on the candidate layers. In our preliminary experi-

ments, this approach improves computational efficiency and slightly enhances model performance. For LLAMA-2-7B, Mistral-7B-v0.1, and LLAMA-3-8B with 32 hidden layers, the candidate layers are set to $\{17,\ldots,32\}$, and only even-numbered layers will be searched. For LLAMA2-13B with 40 hidden layers, the candidate layers are set to $\{31,\ldots,40\}$. All experiments are conducted on a single A100 80GB GPU.

4.1 Analyzing the Distraction Phenomenon in Retrieved Context

We are particularly interested in a challenging QA scenario proposed by (Liu et al., 2024), in which the oracle document is surrounded by numerous semantically similar distractor documents. This configuration challenges the efficacy of the NAIVE LLM-RAG, preventing it from accurately deriving answers based on contextual information, thereby leading to the "Lost in the middle" phenomenon. Specifically, following (Liu et al., 2024), given a query from NaturalQuestions-Open (Kwiatkowski et al., 2019), we select a Wikipedia paragraph containing the answer from the NaturalQuestions annotations as the oracle document. Then, Contriever (Gautier et al., 2022) is employed to extract K-1 additional paragraphs from the Wikipedia corpus that are highly relevant to the query yet do not include the ground truth answer, functioning as distractor documents. The query, the oracle document, and K-1 distractor documents are subsequently processed by the LLM to generate an answer.

From Figure 2, it is evident that the performance of the NAIVE method which concatenates all documents to compose the context, is highly sensitive to the placement of the oracle document within the retrieved documents. Among the four evaluated LLMs, the performance of the NAIVE method significantly deteriorates when the oracle document is neither at the very beginning nor at the end. In contrast, since the proposed LeEns processes each document in parallel during decoding, its performance is naturally independent of the position of the oracle document. In almost all positions, LeEns substantially surpasses the performance of the NAIVE method. Notably, in this challenging scenario, REPLUG which ensemble documents' distributions based on retriever weights perform exceedingly poorly, achieving results merely on par with AvgEns . Based on this observation, we further conduct the weight analysis in Figure 3.




Figure 2: Impact of positioning the oracle document on multi-document question answering performance. A 10-document context typically uses less than 2K tokens; a 20-document context usually uses less than 4K tokens.




Figure 3: (a) The distribution of the similarity difference between the query and the oracle document versus the query and distractor documents. (b) The distribution of the difference in entropy of the first token generated by LLMs when given the oracle document versus distractor documents.

As depicted in Figure 3a, in only approximately 57% of instances, Contriever identifies the oracle document as more similar to the query than the distractor documents. Figure 3b shows the distribution of entropy differences for the first token when conditioned on oracle documents versus distractor documents. It indicates that the entropy is generally lower when the response to a query is based on the oracle document rather than the distractor documents.

Moreover, the proposed CLeHe derived from LeEns further enhances the accuracy of responses in LLMs, particularly evident in the LLAMA2-7B and LLAMA-2-13B. This indicates that contrasting low-entropy distributions of contextual knowledge with high-entropy distributions of parameterized knowledge can further strengthen LLMs' understanding of context.

4.2 Open-Domain Question Answering

Datasets and Metrics. We evaluate our proposed method using four open-domain QA datasets. Natural Questions (Kwiatkowski et al., 2019), Triv-

iaQA (Joshi et al., 2017), WebQ (Berant et al., 2013) and PopQA (Mallen et al., 2023). Natural Questions includes real anonymized queries from Google's search engine. We utilize a filtered test set (Lee et al., 2019) of 3,610 samples with answers limited to no more than five tokens. TriviaQA comprises trivia question-answer pairs that were scraped from the web. We evaluate its development set containing 7,993 samples. WebQ consists of questions generated through the Google Suggest API, with answers that are entities in Freebase. We use its test set of 2,032 samples for evaluation. PopQA is a novel entity-centric opendomain QA dataset that spans a wide range of entity popularity, emphasizing long-tail knowledge. We utilize its test set which includes 14,267 samples for evaluation. For each dataset, we retain only the questions and their corresponding answers. DPR (Karpukhin et al., 2020) is employed to retrieve the top-k passages from the Wikipedia corpus (Dec. 20, 2018) via as evidence documents for each question. Specifically, we report the performance of different decoding methods when retrieving the

Dataset	NQ		TQA			WebQ		PopQA		Gain			
# of Docs	5	10	20	5	10	20	5	10	20	5	10	20	Guin
LLAMA-2-7B													
NAIVE	19.56	26.76	26.37	52.53	59.05	66.00	16.04	18.06	18.35	23.64	25.50	26.10	0.00
AvgEns	14.93	13.51	12.78	51.04	49.84	48.87	12.50	12.06	11.35	14.03	12.49	11.74	-9.40
REPLUG	23.85	23.87	23.92	55.56	55.37	55.31	16.63	16.39	16.77	23.85	23.80	23.91	-1.56
LeEns	25.48	25.79	25.90	61.87	62.72	63.16	20.28	21.65	21.66	27.84	26.57	25.07	+2.50
CLeHe	37.62	36.29	35.48	69.56	69.92	69.72	36.12	36.22	35.77	32.12	31.11	28.89	+11.74
LLAMA-2-13I	В												
Naive	37.98	39.67	29.07	66.82	68.15	69.66	29.83	37.02	27.58	34.77	35.14	31.62	0.00
AvgEns	23.52	20.19	18.71	66.92	65.48	63.86	25.89	24.16	22.64	24.28	22.58	22.02	-8.92
REPLUG	34.12	33.82	34.01	67.83	67.77	67.60	31.25	31.30	30.88	30.79	30.95	31.00	-1.33
LeEns	36.54	34.65	33.63	71.87	72.07	72.16	36.07	35.63	35.42	34.63	33.31	31.72	+1.70
CLeHe	37.31	36.01	34.95	72.24	72.82	72.55	38.19	37.45	36.66	34.22	33.16	31.79	+3.18
Mistral-7B-v0	Mistral-7B-v0.1												
NAIVE	46.20	44.43	42.20	73.53	70.31	73.89	47.69	45.03	40.61	40.23	37.72	38.94	-0.00
AvgEns	40.91	39.36	38.22	76.28	75.73	74.91	47.98	48.13	47.83	37.11	34.61	33.44	-0.52
REPLUG	44.35	44.44	44.58	74.62	74.80	74.60	47.59	47.49	47.21	37.76	37.75	37.78	+1.02
LeEns	46.40	46.45	44.65	78.26	78.97	79.25	49.21	49.70	50.32	42.12	43.31	43.87	+4.32
CLeHe	46.32	46.07	44.64	78.19	78.88	79.14	49.06	49.76	50.37	42.11	43.34	43.89	+4.25
LLAMA-3-8B													
Naive	52.35	51.69	52.33	79.92	81.11	82.10	50.49	50.15	50.12	40.92	42.24	42.95	-0.00
AvgEns	47.12	45.70	44.51	81.31	80.73	79.78	51.82	51.24	51.01	38.95	36.49	35.07	-2.72
REPLUG	50.39	50.33	50.50	79.07	79.16	78.76	50.20	50.79	50.27	39.08	39.24	39.04	-1.63
LeEns	51.74	50.53	49.47	81.80	82.68	83.02	52.17	50.98	51.80	43.63	44.92	45.62	+1.00
CLeHe	52.02	50.78	49.67	81.78	82.84	83.14	51.67	51.62	52.39	43.86	44.84	45.57	+1.17

Table 1: Performance (%) comparison of different decoding methods on benchmark datasets. "Gain" refers to the average absolute improvement (%) across all datasets and different numbers of retrieved documents when compared to the naive baseline.

top-5, top-10, and top-20 documents. Exact Match (EM) metric is utilized for performance evaluation.

4.2.1 Overall Performance

Table 1 presents the overall performance comparison between our proposed method and existing baselines on public benchmark datasets. The results show that when compared to the NAIVE method, our entropy-ensemble-based LeEns demonstrates significant average performance improvements across various LLMs, indicating its superior ability to extract useful information from the context. Moreover, LeEns outperforms REPLUG and AvgEns in almost all settings, indicating that using the uncertainty of LLM output distributions for document scoring more effectively facilitates generating answers than static retriever similarity and unweighted averaging. Comparing LeEns with CLeHe, we observe that further contrasting the ensemble-based low-entropy contextual distribution with the high-entropy distribution of the parametric knowledge leads to performance improvements, particularly noticeable in LLAMA2-7B

and LLAMA2-13B. These observations substantiate that the proposed entropy-based decoding mechanism markedly augments the extraction and utilization of contextual information. Further, on Mistral-7b-v0.1 and LLAMA-3-8B, CLeHe performs similarly to LeEns, indicating no significant enhancement from the contrastive ensemble. We speculate that these two models are less distracted by parametric knowledge when generating answers.

4.2.2 Comparative Analysis with CAD

Our proposed CLeHe shares partial similarities with CAD, notably both employing contrastive decoding to suppress parametric knowledge. The CAD method denotes contrasting the distribution derived from NAIVE LLM-RAG when considering retrieved documents, against the last-layer distribution when the context is not considered. Table 2 shows the performance comparison between our CLeHe and the CAD. It's evident that across most datasets and LLMs, the proposed CLeHe significantly outperforms the CAD, indicating that it is more effective at understanding context.

Model	Method	NQ	TQA	WebQ
LLAMA-2-7B	CAD CLeHe (Ours)	38.92 37.62	65.92 69.56	30.77 36.12
LLAMA-2-13B	CAD CLeHe (Ours)	39.43 37.31	68.96 70.12	33.76 38.19
Mistral-7B-v0.1	CAD CLeHe (Ours)	44.32 46.32	70.27 78.26	41.49 49.76
LLAMA-3-8B	CAD CLeHe (Ours)	51.80 52.02	79.01 81.78	46.26 51.62

Table 2: Performance under the top-5 document setting.




Figure 4: Effect of layer selection strategies on WebQ under the top-5 document setting.

4.2.3 Analysis of Layer Selection Strategies

In our CLeHe, we dynamically select the layer with the highest entropy distribution (Eq.(6)) as the representative for parameterized knowledge at each generation time step. A simpler alternative method involves using the last layer to represent the parametric knowledge. A more complete strategy is based on distribution differences (Chuang et al., 2023) which has been demonstrated to enhance the factuality of LLMs. For this strategy, we first calculate the JSD (Jensen-Shannon Divergence) between the low-entropy ensemble distribution (Eq.(3)) and the distribution of each layer when without context. Then the distribution with the highest JSD value for contrast is selected for contrast. As shown in Figure 4, our proposed entropybased layer selection strategy surpasses the fixed final-layer strategy. Also, an intriguing observation is that the performance of this entropy-based dy-




Figure 5: Hyper-parameter analysis using 1K evaluation samples of NQ under the top-5 document setting.

namic selection strategy slightly exceeds the best performance achievable by fixed static layers. Appendix A.1 visualizes the average entropy of tokens generated when using different layers for contrast, which shows that layers with higher entropy often correspond to better performance. Meanwhile, it performs better overall than the dynamic JSD-based strategy. It's noteworthy that in our setting, using the JSD-based selection strategy for contrast is unstable and results in considerable performance degradation in LLAMA-2-13B.

4.2.4 Hyper-Parameter Analysis

We study the influence of the introduced hyperparameters: τ and β . As shown in Figure 5, a small value of τ (*e.g.*, 0.1 or 0.25) typically results in better performance; as τ increases, the performance gradually declines. Ideally, when $\tau \to \infty$, the performance of the proposed LeEns will match that of AvgEns. Regarding β , it's observed that for LLAMA-2-7B, a high β (e.g., 5) enables it to effectively contrast the differences between external and parametric knowledge for improved performance. For other evaluated LLMs, we suggest setting it to a small value, saying [0.25, 0.5].

5 Conclusions

In this paper, we proposed a novel decoding method that is guided by entropy considerations to mitigate the distractibility issue from both external retrieved documents and parametric knowledge. First, we conducted parallel retrieval augmentation with entropy-based ensemble weight to obtain the low-entropy distribution of context. Furthermore, we contrasted this distribution against the highest-entropy distribution among layers when without context to amplify the external knowledge preserved in context. Extensive experiments showed the proposed method's effectiveness in retrieval-augmented open-domain question answering.

6 Limitations

One limitation of our study is that we only validated the effectiveness of our method on question answering datasets, without testing it on other knowledge-intensive tasks such as fact verification. Extending the method proposed in this paper to other retrieval-augmented scenarios will be a future research direction.

Additionally, due to computational power constraints, we only tested the effectiveness of the proposed method on models with fewer than 13B parameters. However, whether the method proposed in this paper is applicable to LLMs with more parameters (*e.g.*, 70B or more) remains to be explored in future research.

References

- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. 2013. Semantic parsing on freebase from question-answer pairs. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1533–1544.
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George Bm Van Den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, et al. 2022. Improving language models by retrieving from trillions of tokens. In *International conference on machine learning*, pages 2206–2240. PMLR.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. 2023. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*.
- Yung-Sung Chuang, Yujia Xie, Hongyin Luo, Yoon Kim, James R Glass, and Pengcheng He. 2023. Dola: Decoding by contrasting layers improves factuality in large language models. In *The Twelfth International Conference on Learning Representations*.
- Izacard Gautier, Caron Mathilde, Hosseini Lucas, Riedel Sebastian, Bojanowski Piotr, Joulin Armand, and Grave Edouard. 2022. Unsupervised dense information retrieval with contrastive learning. *Transactions on Machine Learning Research*.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Mingwei Chang. 2020. Retrieval augmented

- language model pre-training. In *International conference on machine learning*, pages 3929–3938. PMLR.
- Geoffrey E Hinton. 2002. Training products of experts by minimizing contrastive divergence. *Neural computation*, 14(8):1771–1800.
- Gautier Izacard and Edouard Grave. 2020. Leveraging passage retrieval with generative models for open domain question answering. *arXiv* preprint *arXiv*:2007.01282.
- Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. 2023. Atlas: Few-shot learning with retrieval augmented language models. *Journal of Machine Learning Research*, 24(251):1–43.
- Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. 2017. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. In *ACL* (1), pages 1601–1611. Association for Computational Linguistics.
- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick S. H. Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *EMNLP* (1), pages 6769–6781. Association for Computational Linguistics.
- Jungo Kasai, Keisuke Sakaguchi, Ronan Le Bras, Akari Asai, Xinyan Yu, Dragomir Radev, Noah A Smith, Yejin Choi, Kentaro Inui, et al. 2024. Realtime qa: What's the answer right now? *Advances in Neural Information Processing Systems*, 36.
- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, et al. 2019. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466.
- Kenton Lee, Ming-Wei Chang, and Kristina Toutanova. 2019. Latent retrieval for weakly supervised open domain question answering. In *ACL* (1), pages 6086–6096. Association for Computational Linguistics.
- Sicong Leng, Hang Zhang, Guanzheng Chen, Xin Li, Shijian Lu, Chunyan Miao, and Lidong Bing. 2023. Mitigating object hallucinations in large vision-language models through visual contrastive decoding. arXiv preprint arXiv:2311.16922.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474.

- Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. 2016. A diversity-promoting objective function for neural conversation models. In *HLT-NAACL*, pages 110–119. The Association for Computational Linguistics.
- Xiang Lisa Li, Ari Holtzman, Daniel Fried, Percy Liang, Jason Eisner, Tatsunori Hashimoto, Luke Zettlemoyer, and Mike Lewis. 2022. Contrastive decoding: Open-ended text generation as optimization. *arXiv* preprint arXiv:2210.15097.
- Xi Victoria Lin, Xilun Chen, Mingda Chen, Weijia Shi, Maria Lomeli, Rich James, Pedro Rodriguez, Jacob Kahn, Gergely Szilvasy, Mike Lewis, et al. 2023. Ra-dit: Retrieval-augmented dual instruction tuning. *arXiv preprint arXiv:2310.01352*.
- Alisa Liu, Maarten Sap, Ximing Lu, Swabha Swayamdipta, Chandra Bhagavatula, Noah A. Smith, and Yejin Choi. 2021. Dexperts: Decoding-time controlled text generation with experts and anti-experts. In *ACL/IJCNLP* (1), pages 6691–6706. Association for Computational Linguistics.
- Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the middle: How language models use long contexts. *Transactions of the Association* for Computational Linguistics, 12:157–173.
- Shayne Longpre, Kartik Perisetla, Anthony Chen, Nikhil Ramesh, Chris DuBois, and Sameer Singh. 2021. Entity-based knowledge conflicts in question answering. *arXiv preprint arXiv:2109.05052*.
- Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi. 2023. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In *ACL* (1), pages 9802–9822. Association for Computational Linguistics.
- Sean O'Brien and Mike Lewis. 2023. Contrastive decoding improves reasoning in large language models. *arXiv preprint arXiv:2309.09117*.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744.
- Fabio Petroni, Aleksandra Piktus, Angela Fan, Patrick Lewis, Majid Yazdani, Nicola De Cao, James Thorne, Yacine Jernite, Vladimir Karpukhin, Jean Maillard, et al. 2020. Kilt: a benchmark for knowledge intensive language tasks. *arXiv preprint arXiv:2009.02252*.
- Zexuan Qiu, Jingjing Li, Shijue Huang, Wanjun Zhong, and Irwin King. 2024. Clongeval: A chinese benchmark for evaluating long-context large language models. *arXiv preprint arXiv:2403.03514*.

- Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. In-context retrieval-augmented language models. *Transactions of the Association for Computational Linguistics*, 11:1316–1331.
- Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H Chi, Nathanael Schärli, and Denny Zhou. 2023a. Large language models can be easily distracted by irrelevant context. In *International Conference on Machine Learning*, pages 31210–31227. PMLR.
- Weijia Shi, Xiaochuang Han, Mike Lewis, Yulia Tsvetkov, Luke Zettlemoyer, and Scott Wen-tau Yih. 2023b. Trusting your evidence: Hallucinate less with context-aware decoding. *arXiv preprint arXiv*:2305.14739.
- Weijia Shi, Sewon Min, Michihiro Yasunaga, Minjoon Seo, Rich James, Mike Lewis, Luke Zettlemoyer, and Wen-tau Yih. 2023c. Replug: Retrieval-augmented black-box language models. *arXiv* preprint arXiv:2301.12652.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- Liam Van der Poel, Ryan Cotterell, and Clara Meister. 2022. Mutual information alleviates hallucinations in abstractive summarization. *arXiv preprint arXiv:2210.13210*.
- Neeraj Varshney, Wenlin Yao, Hongming Zhang, Jianshu Chen, and Dong Yu. 2023. A stitch in time saves nine: Detecting and mitigating hallucinations of llms by validating low-confidence generation. *arXiv* preprint arXiv:2307.03987.
- Jonas Waldendorf, Barry Haddow, and Alexandra Birch. 2024. Contrastive decoding reduces hallucinations in large multilingual machine translation models. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2526–2539.
- Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. 2023. Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts. In *The Twelfth International Conference on Learning Representations*.
- Howard Yen, Tianyu Gao, and Danqi Chen. 2024. Long-context language modeling with parallel context encoding. *arXiv preprint arXiv:2402.16617*.
- Tianjun Zhang, Shishir G Patil, Naman Jain, Sheng Shen, Matei Zaharia, Ion Stoica, and Joseph E Gonzalez. 2024. Raft: Adapting language model to domain specific rag. *arXiv* preprint arXiv:2403.10131.

A Appendix

A.1 Relations between Entropy and Performance

Figure 6 shows the average entropy values of tokens generated on the WebQ dataset by LLAMA2-7B and LLAMA2-13B when using different layers for contrast. For LLAMA2-7B, the entropy initially shows a decreasing trend and later an increasing trend among the tested layers, which correlates somewhat with the overall performance trend. For LLAMA2-13B, it was found that both the highest performance layer and the layer with the highest entropy are the last layer. This might explain why the performance of our entropy maximization-based layer selection is overall consistent with the performance of the last layer (see Figure 4).

A.2 Illustration of the False Positive Case

Figure 7 illustrates a false positive scenario. In this example, the retrieval-augmented model assigns high probabilities to the words "Washington", "New York", and "Columbia" as candidate positives. However, in the low-entropy output of a specific layer (typically the last layer) without context, the probability assigned to "Columbia" is notably low. If contrastive decoding is applied, it would mistakenly increase "Columbia's" probability, leading to an incorrect prediction.

A.3 Prompt Format

```
Write a high-quality answer for the given question using only the provided search results.

Document [1] (Title: ...) ...

Document [2] (Title: ...) ...

Question: {Question}

Answer:

Write a high-quality answer for the given question using only the provided search results.

Document (Title: ...) ...

Question: {Question}

Answer:
```

Table 3: The Prompt format for NAIVE RAG (Top); and the prompt format for the proposed CLeHe (Bottom) to process each document in parallel.

In Table 3 (Bottom), we showcase the prompt format utilized for the proposed CLeHe; its struc-

ture is basically the same as that of the NAIVE RAG. The primary distinction is that due to the parallel processing, we need to repeat the task prompt and question for each document.




Figure 6: Average entropy of generated tokens and performance for each layer.




Figure 7: Illustration of the false positive case.